

```
In [3]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score, confusion_matrix, roc_auc_score, roc_curve
import matplotlib.pyplot as plt
import numpy as np
```

```
In [4]: data = pd.read_csv('heart_failure_clinical_records_dataset.csv')
```

```
In [5]: data.describe()
```

```
Out[5]:
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction
<b>count</b>	299.000000	299.000000	299.000000	299.000000	299.000000
<b>mean</b>	60.833893	0.431438	581.839465	0.418060	38.083612
<b>std</b>	11.894809	0.496107	970.287881	0.494067	11.834841
<b>min</b>	40.000000	0.000000	23.000000	0.000000	14.000000
<b>25%</b>	51.000000	0.000000	116.500000	0.000000	30.000000
<b>50%</b>	60.000000	0.000000	250.000000	0.000000	38.000000
<b>75%</b>	70.000000	1.000000	582.000000	1.000000	45.000000
<b>max</b>	95.000000	1.000000	7861.000000	1.000000	80.000000

```
In [6]: data.isnull().sum()
```

```
Out[6]: age                                0
anaemia                                    0
creatinine_phosphokinase                  0
diabetes                                  0
ejection_fraction                         0
high_blood_pressure                       0
platelets                                 0
serum_creatinine                          0
serum_sodium                             0
sex                                        0
smoking                                   0
time                                      0
DEATH_EVENT                              0
dtype: int64
```

```
In [7]: data.corr()['DEATH_EVENT'].sort_values()
```

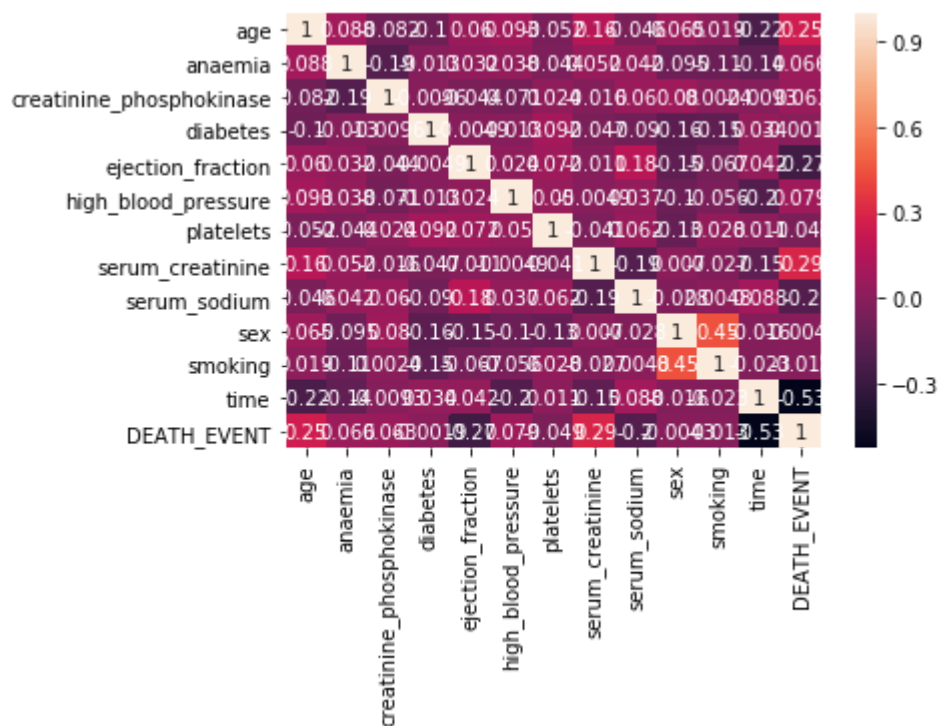
```
Out[7]: time                -0.526964
ejection_fraction          -0.268603
serum_sodium               -0.195204
platelets                  -0.049139
smoking                    -0.012623
sex                        -0.004316
diabetes                   -0.001943
creatinine_phosphokinase   0.062728
anaemia                    0.066270
high_blood_pressure        0.079351
age                        0.253729
serum_creatinine           0.294278
DEATH_EVENT                1.000000
Name: DEATH_EVENT, dtype: float64
```

```
In [8]: plt.figure(figsize=(12, 6))
```

```
Out[8]: <Figure size 864x432 with 0 Axes>
<Figure size 864x432 with 0 Axes>
```

```
In [9]: sns.heatmap(data.corr(), annot=True)
```

```
Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1d7b23a36a0>
```



```
In [10]: X = data.drop('DEATH_EVENT', axis=1)
y = data['DEATH_EVENT']
```

```
In [11]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.8, random_state=29)
```

```
In [12]: rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
In [13]: # Train the classifier using cross-validation
cross_val_scores = cross_val_score(rf_classifier, X_train, y_train, cv=5, scoring='accuracy')
mean_cross_val_accuracy = cross_val_scores.mean()
```

```
In [14]: # Train the classifier on the full training set
rf_classifier.fit(X_train, y_train)
```

```
Out[14]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=1,
                                oob_score=False, random_state=42, verbose=0, warm_start=False)
```

```
In [15]: # Make predictions on the test set
predictions = rf_classifier.predict(X_test)
```

```
In [17]: # Evaluate and print the accuracy
accuracy = accuracy_score(y_test, predictions)
print(f"Accuracy (Before Cross-Validation): {accuracy:.4f}")
```

Accuracy (Before Cross-Validation): 0.8500

```
In [18]: # Display Cross-Validation Mean Accuracy
print(f"Cross-Validation Mean Accuracy: {mean_cross_val_accuracy:.4f}")
```

Cross-Validation Mean Accuracy: 0.8326

```
In [19]: # Confusion Matrix
conf_matrix = confusion_matrix(y_test, predictions)
print("Confusion Matrix:")
print(conf_matrix)
```

Confusion Matrix:

```
[[144  23]
 [ 13  60]]
```

```
In [20]: # AUC-ROC Curve
y_proba = rf_classifier.predict_proba(X_test)[:, 1]
roc_auc = roc_auc_score(y_test, y_proba)
fpr, tpr, _ = roc_curve(y_test, y_proba)
```


```
In [21]: plt.figure(figsize=(8, 6))
```

Out[21]: <Figure size 576x432 with 0 Axes>

<Figure size 576x432 with 0 Axes>


```
In [22]: plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (AUC = {:.2f})'.format(roc_auc))
```

```
Out[22]: [<matplotlib.lines.Line2D at 0x1d7b276a8d0>]
```




```
In [23]: plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
```

```
Out[23]: [<matplotlib.lines.Line2D at 0x1d7b27c60b8>]
```




```
In [24]: plt.xlabel('False Positive Rate')
```

```
Out[24]: Text(0.5,0,'False Positive Rate')
```



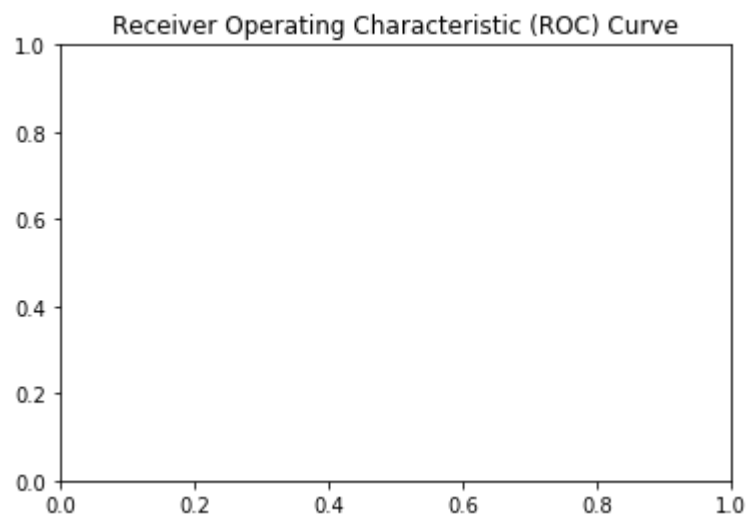
```
In [25]: plt.ylabel('True Positive Rate')
```

```
Out[25]: Text(0,0.5,'True Positive Rate')
```




```
In [26]: plt.title('Receiver Operating Characteristic (ROC) Curve')
```

```
Out[26]: Text(0.5,1,'Receiver Operating Characteristic (ROC) Curve')
```



```
In [29]: plt.legend(loc='lower right')  
plt.show()
```

No handles with labels found to put in legend.



```
In [32]: from sklearn.metrics import log_loss
```

```
In [33]: # Accuracy and Log Loss Curves during Training
train_acc = []
test_acc = []
train_logloss = []
test_logloss = []
```

```
In [34]: for n_estimators in range(1, 101):
    rf_classifier = RandomForestClassifier(n_estimators=n_estimators, random_state=42)
    rf_classifier.fit(X_train, y_train)

    # Training Accuracy and Log Loss
    train_preds = rf_classifier.predict(X_train)
    train_accuracy = accuracy_score(y_train, train_preds)
    train_log_loss = log_loss(y_train, rf_classifier.predict_proba(X_train))

    # Test Accuracy and Log Loss
    test_preds = rf_classifier.predict(X_test)
    test_accuracy = accuracy_score(y_test, test_preds)
    test_log_loss = log_loss(y_test, rf_classifier.predict_proba(X_test))

    # Append values
    train_acc.append(train_accuracy)
    test_acc.append(test_accuracy)
    train_logloss.append(train_log_loss)
    test_logloss.append(test_log_loss)
```


```
In [35]: # Plotting Accuracy Curves
plt.figure(figsize=(12, 4))
```

```
Out[35]: <Figure size 864x288 with 0 Axes>
```

```
<Figure size 864x288 with 0 Axes>
```


```
In [36]: plt.subplot(1, 2, 1)
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x1d7b29a9780>
```




```
In [37]: plt.plot(range(1, 101), train_acc, label='Training Accuracy')
```

```
Out[37]: [ <matplotlib.lines.Line2D at 0x1d7b2a1cdd8>]
```




```
In [38]: plt.plot(range(1, 101), test_acc, label='Testing Accuracy')
```

```
Out[38]: [ <matplotlib.lines.Line2D at 0x1d7b2a7b908>]
```




```
In [39]: plt.xlabel('Number of Trees (n_estimators)')
```

```
Out[39]: Text(0.5,0,'Number of Trees (n_estimators)')
```



```
In [40]: plt.ylabel('Accuracy')
```


```
Out[40]: Text(0,0.5,'Accuracy')
```



```
In [41]: plt.legend()
```


No handles with labels found to put in legend.

```
Out[41]: <matplotlib.legend.Legend at 0x1d7b2b592b0>
```



```
In [42]: # Plotting Log Loss Curves  
plt.subplot(1, 2, 2)
```


```
Out[42]: <matplotlib.axes._subplots.AxesSubplot at 0x1d7b3b6fa58>
```






```
In [43]: plt.plot(range(1, 101), train_logloss, label='Training Log Loss')
```

```
Out[43]: [ <matplotlib.lines.Line2D at 0x1d7b3be2908>]
```




```
In [44]: plt.plot(range(1, 101), test_logloss, label='Testing Log Loss')
```

```
Out[44]: [ <matplotlib.lines.Line2D at 0x1d7b3c3b710>]
```




```
In [45]: plt.xlabel('Number of Trees (n_estimators)')
```

```
Out[45]: Text(0.5,0,'Number of Trees (n_estimators)')
```



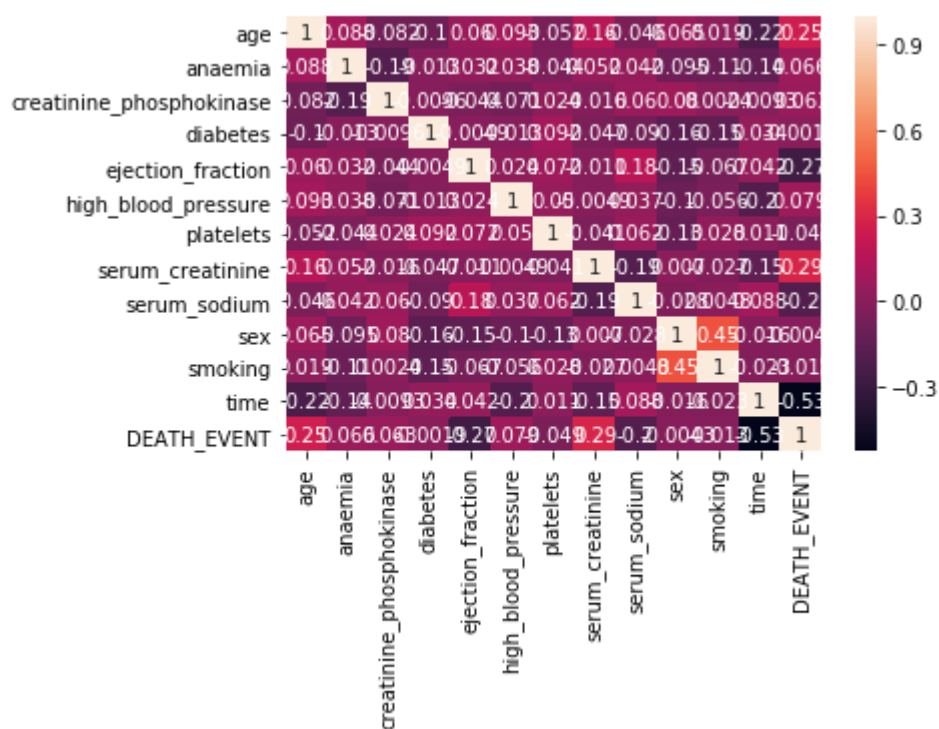
```
In [46]: plt.ylabel('Log Loss')
```

```
Out[46]: Text(0,0.5,'Log Loss')
```



```
In [52]: sns.heatmap(data.corr(), annot=True)
```

```
Out[52]: <matplotlib.axes._subplots.AxesSubplot at 0x1d7b3c14eb8>
```



```
In [53]: accuracy = accuracy_score(y_test, predictions)
print(f"Accuracy (Before Cross-Validation): {accuracy:.4f}")
```

Accuracy (Before Cross-Validation): 0.8500

```
In [54]: # Display Cross-Validation Mean Accuracy
print(f"Cross-Validation Mean Accuracy: {mean_cross_val_accuracy:.4f}")
```

Cross-Validation Mean Accuracy: 0.8326